

Outlines

- 0. Tensor basics
 - How to create and manipulate tensors in PyTorch as well as how to understand tensors as a particular data structure/type
- 1. Simple Linear Regression
 - General intuition about machine learning
 - Standard way of building and training models
- 2. Feedforward neural network
 - Non-linear function
 - Regression → classification
 - Training + testing
- 3. Text classification
 - When the input data are not numerical: text data transformation
 - Application of word embedding, BOW, LSTM
 - More realistic workflow: batching, training + testing + deploying etc.

How a gradient-based machine learning model typically works

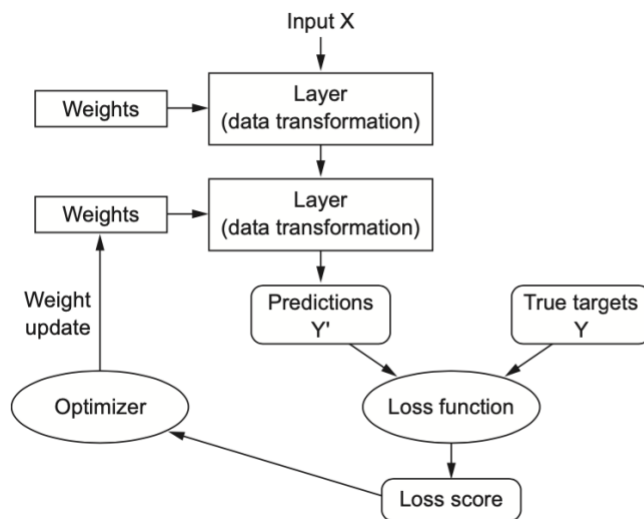


Figure 3.1 Relationship between the network, layers, loss function, and optimizer

Credit: Deep Learning with Python (2nd), François Chollet, 2022, p58